

# Reinforcement Learning: A Brief Introduction

*Yiwen Song*

[gavinsyw@sjtu.edu.cn](mailto:gavinsyw@sjtu.edu.cn)

Reference Book:

**Machine Learning** by Zhihua Zhou, Tsinghua University Press, Chap. 6, 16

**Reinforcement Learning: An Introduction** by Richard S. Sutton and Andrew G. Barto, MIT Press, Chap. 1-6, 9, 16, 17

# Contents

- Markov Decision Process: the Environment
- Traditional RL: TD, Q-, and SARSA
- Convergence: Whether & How fast
- Continuous State: Value Function Approximation
- Deep RL: Pros and Cons
- Industrial Examples

# Markov Decision Process: the Environment


[Wikipedia, 2019] [https://en.wikipedia.org/wiki/Reinforcement\\_learning](https://en.wikipedia.org/wiki/Reinforcement_learning).

[Wikipedia, 2019] [https://en.wikipedia.org/wiki/Markov\\_decision\\_process](https://en.wikipedia.org/wiki/Markov_decision_process)

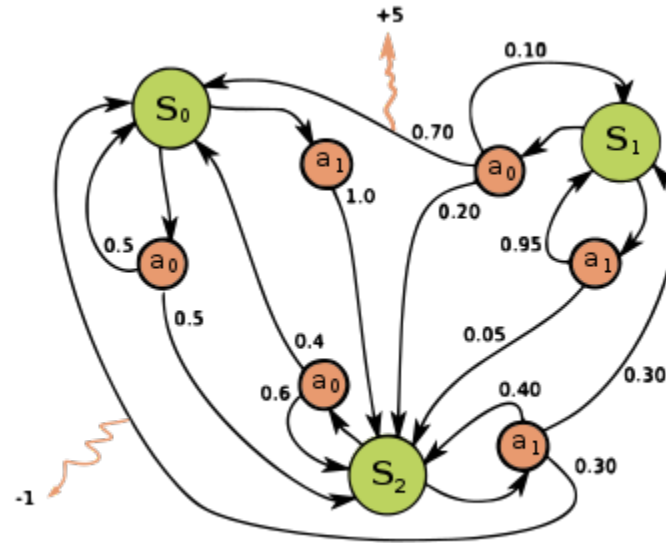
[Z. Zhou, 2016] Machine Learning, Tsinghua University Press

[S. Richard, et al., 2018] Reinforcement Learning: An Introduction, MIT Press

# The interaction of an agent and the environment forms the basic model of RL.



The Markov Decision Process is a 4(5) tuple  
 $(\mathcal{S}, \mathcal{A}, P_a, R_a, \gamma)$



# There are 2 traditional solutions for MDP: Policy Iteration and Value Iteration

- $V(s) := \sum_{s'} P_{\pi(s)}(s, s') (R_{\pi(s)}(s, s') + \gamma V(s'))$
- $\pi(s) := \arg \max_a [\sum_{s'} P(s'|s, a) (R(s'|s, a) + \gamma V(s'))]$
- OR
- $V_{i+1}(s) := \max_a [\sum_{s'} P_a(s, s') (R_a(s, s') + \gamma V(s'))]$

# Traditional RL: MC, TD, Q-, and SARSA

[Wikipedia, 2019] [https://en.wikipedia.org/wiki/Temporal\\_difference\\_learning](https://en.wikipedia.org/wiki/Temporal_difference_learning)

[Wikipedia, 2019] <https://en.wikipedia.org/wiki/Q-learning>

[Wikipedia, 2019] <https://en.wikipedia.org/wiki/State-action-reward-state-action>

[Z. Zhou, 2016] Machine Learning, Tsinghua University Press

[S. Richard, et al., 2018] Reinforcement Learning: An Introduction, MIT Press

[R. Sutton, et al., 1995] Temporal difference learning and TD-Gammon, *Communications of the ACM*

[C. Watkins, et al., 1992] Q-learning, *Machine Learning*

[G. Rummery, et al., 1994] On-line Q-learning using connectionist systems, *Technical Report CUED/F-INFENG/TR*

# Monte-Carlo Method solves the problem without complete knowledge of environment

- Input: environment  $E$ , action space  $A$ , start state  $x_0$ , total step  $T$
- Output: policy  $\pi$
- 2 steps: 1. Estimating  $v_\pi$ ; 2. Gaining  $\pi^*$
- $V(x_t, a_t) = \frac{V(x_t, a_t) \times n + R}{n+1}$ ,  $R = \frac{1}{T-t} \sum_{i=t+1}^T r_i$
- $\epsilon$ -greedy policy for evaluating policy:
$$\pi(x) = \begin{cases} \arg \max_{a'} V(x, a'), & \text{with probability } 1 - \epsilon \\ \text{randomly choose an action from } A, & \text{with probability } \epsilon \end{cases}$$



## First-visit MC prediction, for estimating $V \approx v_\pi$

Input: a policy  $\pi$  to be evaluated

Initialize:

$V(s) \in \mathbb{R}$ , arbitrarily, for all  $s \in \mathcal{S}$

$Returns(s) \leftarrow$  an empty list, for all  $s \in \mathcal{S}$

Loop forever (for each episode):

Generate an episode following  $\pi$ :  $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :

$G \leftarrow \gamma G + R_{t+1}$

Unless  $S_t$  appears in  $S_0, S_1, \dots, S_{t-1}$ :

Append  $G$  to  $Returns(S_t)$

$V(S_t) \leftarrow \text{average}(Returns(S_t))$

## On-policy first-visit MC control (for $\varepsilon$ -soft policies), estimates $\pi \approx \pi_*$

Algorithm parameter: small  $\varepsilon > 0$

Initialize:

$\pi \leftarrow$  an arbitrary  $\varepsilon$ -soft policy

$Q(s, a) \in \mathbb{R}$  (arbitrarily), for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$

$Returns(s, a) \leftarrow$  empty list, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$

Repeat forever (for each episode):

Generate an episode following  $\pi$ :  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :

$G \leftarrow \gamma G + R_{t+1}$

Unless the pair  $S_t, A_t$  appears in  $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$ :

Append  $G$  to  $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{average}(Returns(S_t, A_t))$

$A^* \leftarrow \operatorname{argmax}_a Q(S_t, a)$  (with ties broken arbitrarily)

For all  $a \in \mathcal{A}(S_t)$ :

$$\pi(a|S_t) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(S_t)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(S_t)| & \text{if } a \neq A^* \end{cases}$$

# Temporal Difference (TD) Learning is an advance of Monte-Carlo Method **applying DP**

- Apply DP in the process of sampling (experiencing). Monte-Carlo Algorithm don't care what happened before

$$V(s_t) \leftarrow V(s_t) + \alpha[G_t - V(s_t)]$$

- However, TD consider both the observed reward and the estimate

$$V(s_t) \leftarrow V(s_t) + \alpha[R_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

- Thus the estimating procedure goes faster

## Tabular TD(0) for estimating $v_\pi$

Input: the policy  $\pi$  to be evaluated

Algorithm parameter: step size  $\alpha \in (0, 1]$

Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$

Loop for each episode:

    Initialize  $S$

    Loop for each step of episode:

$A \leftarrow$  action given by  $\pi$  for  $S$

        Take action  $A$ , observe  $R, S'$

$V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$

$S \leftarrow S'$


    until  $S$  is terminal

# The example **Random Walk** shows how TD performs better than MC




- A random walk is a **Markov Reward Process (MRP)**.
- Starting point: C; move left or right with same probability; terminates on extreme left or right.
- True value:  $v_{\pi}(A \sim E) = \frac{1}{6} \sim \frac{5}{6}$

## Result of TD(0), with different episodes



## Comparison between TD(0) and MC, starting at point C



# Q-Learning: Off-policy TD Control

## Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+$ ,  $a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Loop for each step of episode:

        Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

        Take action  $A$ , observe  $R, S'$

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

    until  $S$  is terminal

# State-Action-Reward-State-Action (SARSA): On-policy TD-Control

**Sarsa (on-policy TD control) for estimating  $Q \approx q_*$**

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

    Loop for each step of episode:

        Take action  $A$ , observe  $R, S'$

        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

    until  $S$  is terminal



# Convergence: Whether and How fast?

[S. Singh, et al., 2000] Convergence Results for Single-Step On-Policy Reinforcement-Learning Algorithms, *Machine Learning*

[V. Borkar, et al., 2000] The O.D.E. Method for Convergence of Stochastic Approximation and Reinforcement Learning, *SIAM Control Optimization*

[W. Beggs, 2005] On the convergence of reinforcement learning, *Journal of Economic Theory*

[C. Szepevari, et al., 1999] A unified analysis of value-function-based reinforcement-learning algorithms, *Neural Computation*

[S. Russel, et al., 1999] Convergence of reinforcement learning with general function approximators, *IJCAI*

[J. Tsitsiklis, 2003] On the Convergence of Optimistic Policy Iteration, *The Journal of Machine Learning Research*

# Convergence of Reinforcement Learning

<http://www.cs.cmu.edu/afs/cs.cmu.edu/project/learn-43/lib/idauction2/.g/web/glossary/converge.html>

		Not Residual			Residual		
		Fixed distribution(On-policy)	Fixed Distribution	Usually-greedy distribution	Fixed distribution(On-policy)	Fixed Distribution	Usually-greedy distribution
Markov Chain	Lookup table	Y	Y		Y	Y	
	Averager	Y	Y		Y	Y	
	Linear	Y	N		Y	Y	
	Non-linear	N	N		Y	Y	
MDP	Lookup table	Y	Y	Y	Y	Y	Y
	Averager	Y	Y	N	Y	Y	Y
	Linear	N	N	N	Y	Y	Y
	Non-linear	N	N	N	Y	Y	Y
POMDP	Lookup table	N	N	N	Y	Y	Y
	Averager	N	N	N	Y	Y	Y
	Linear	N	N	N	Y	Y	Y
	Non-linear	N	N	N	Y	Y	Y

# We should assure convergence when applying RL algorithms. However,

- **Convergence may be very slow when the state space is huge,** which is a major weakness of reinforcement learning.
- Two challenges:
  - No general theorem to infer convergence of a specific RL model.
  - Convergence difficult to infer for
    - **Non-i.i.d. samples**
    - **Dynamically changing learning policy (Usually in DRL)**

There are some tricks, however, to assure convergence speed, sacrificing convergence.

- For example:
- A simple TD-RL model (Q-, SARSA) converges as  $t \rightarrow \infty$ .
- The convergence ratio is  $\frac{c_1}{t} + c_2 \log(1 + \frac{c_3}{t}) \rightarrow 0$
- By some means we can achieve a convergence ratio  $c_4 e^{-c_5 t} + c_6$ , sacrificing convergence for learning speed.

[S. Zou, 2019] Non-asymptotic Analysis of Reinforcement Learning Algorithms with Function Approximation (Seminar)

# Continuous State: Value Function Approximation

[Z. Zhou, 2016] Machine Learning, Tsinghua University Press

[S. Richard, et al., 2018] Reinforcement Learning: An Introduction, MIT Press

[L. Busoniu, et al., 2010] Reinforcement Learning Dynamic Programming Using Function Approximation, Hall/CRC Press

# Initiative of Value Function Approximation

- Consider an MDP with continuous state space.
- Can we just discretize the state space?
- Sometimes, yes. But not always yes.
- Therefore, why not just use a function to represent the value function?

$$V(s): \mathbb{R}^n \rightarrow \mathbb{R}, \quad S: \mathbb{R}^n \rightarrow \mathbb{R}^n$$

# The simplest case of VFA is **Linear-VFA**

- Take state space  $X = \mathbb{R}^n$ , and let  $V_{\boldsymbol{\theta}}(\mathbf{x}) = \boldsymbol{\theta}^T \mathbf{x}$
- We want to minimize the RMS error of our approximation:

$$E_{\boldsymbol{\theta}} = \mathbb{E}_{\mathbf{x} \sim \pi} \left[ \left( V^{\pi}(\mathbf{x}) - V_{\boldsymbol{\theta}}(\mathbf{x}) \right)^2 \right]$$

- Apply gradient descent and we have

$$\boldsymbol{\theta} = \boldsymbol{\theta} + \alpha \left( V^{\pi}(\mathbf{x}) - V_{\boldsymbol{\theta}}(\mathbf{x}) \right) \mathbf{x}$$

- Apply TD-Learning and we have the update method for  $\boldsymbol{\theta}$

$$\boldsymbol{\theta} = \boldsymbol{\theta} + \alpha (r + \gamma \boldsymbol{\theta}^T \mathbf{x}' - \boldsymbol{\theta}^T \mathbf{x}) \mathbf{x}$$

### Semi-gradient TD(0) for estimating $\hat{v} \approx v_\pi$

Input: the policy  $\pi$  to be evaluated

Input: a differentiable function  $\hat{v} : \mathcal{S}^+ \times \mathbb{R}^d \rightarrow \mathbb{R}$  such that  $\hat{v}(\text{terminal}, \cdot) = 0$

Algorithm parameter: step size  $\alpha > 0$

Initialize value-function weights  $\mathbf{w} \in \mathbb{R}^d$  arbitrarily (e.g.,  $\mathbf{w} = \mathbf{0}$ )

Loop for each episode:

    Initialize  $S$

    Loop for each step of episode:

        Choose  $A \sim \pi(\cdot|S)$

        Take action  $A$ , observe  $R, S'$

$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})] \nabla \hat{v}(S, \mathbf{w})$

$S \leftarrow S'$

    until  $S$  is terminal



# More complex models work, such as **Kernel-VFA**

- Use **Kernel Function** as approximation for value function.
- Common Kernel Functions:

Name	Expression	Parameters
Linear Kernel	$\mathcal{K}(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$	
Polynomial Kernel	$\mathcal{K}(\mathbf{x}_i, \mathbf{x}_j) = (\mathbf{x}_i^T \mathbf{x}_j)^d$	$d \geq 1$
Gaussian Kernel	$\mathcal{K}(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{\ \mathbf{x}_i - \mathbf{x}_j\ ^2}{2\sigma^2}\right)$	$\sigma > 0$
Laplace Kernel	$\mathcal{K}(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{\ \mathbf{x}_i - \mathbf{x}_j\ }{\sigma}\right)$	$\sigma > 0$
Sigmoid Kernel	$\mathcal{K}(\mathbf{x}_i, \mathbf{x}_j) = \tanh(\beta \mathbf{x}_i^T \mathbf{x}_j + \theta)$	$\beta > 0, \theta < 0$

# Deep RL: Pros and Cons

[P. Dimitri, 1996] Neuro-Dynamic Programming, MIT Press

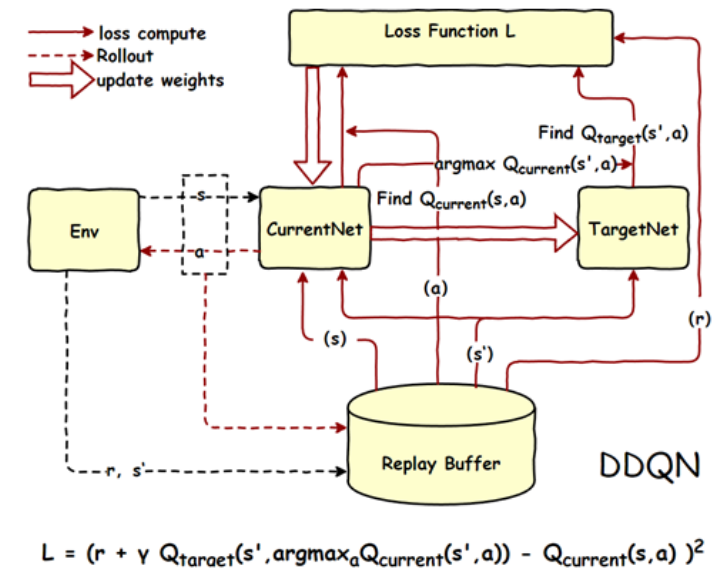
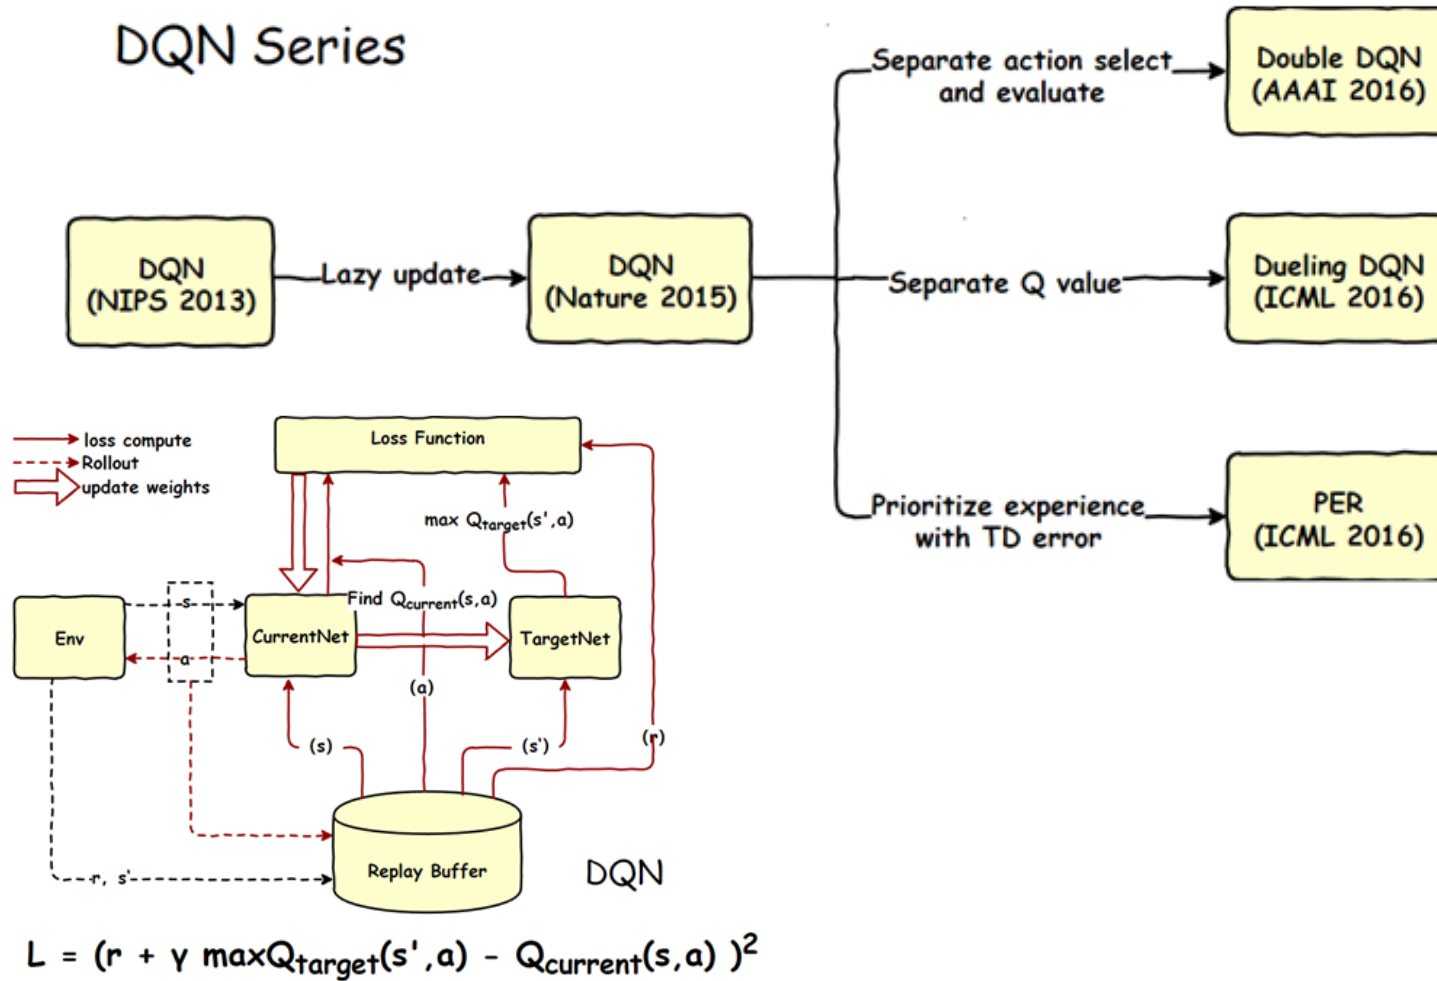
[V. Mnih, et al., 2013] Playing Atari with Deep Reinforcement Learning, *NIPS*

[V. Mnih, et al., 2015] Human-level control through deep reinforcement learning, *Nature*

[V. Hasselt, et al, 2015] Deep Reinforcement Learning with Double Q-learning, *AAAI*

# Revolution of DQN

## DQN Series



# Advantage of DQN

- Compared to Traditional Q
  - Multi-sampling during a step
  - Random sampling increases efficiency for continuous data
  - Avoid learning divergence by experiencing playback
- Compared to other Neural Networks
  - Better Interpretability
  - Can be manually corrected by **Apprenticeship Learning**

# Disadvantage of DQN

- Uncertainty of convergence
- Great cost in hardware
- Sensitive to disturbance during training

# Industrial Examples: When do we use RL?

[S. Richard, et al., 2018] Reinforcement Learning: An Introduction, MIT Press

[S. David, et al., 2018] A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play, *Science*

# Apprenticeship Learning is usually applied in Industrial RL.

- Two ways:
  - **Direct Apprenticeship Learning:** Learn from an expert's action first, then use on-policy training to get the policy;
  - **Inverse Reinforcement Learning:** discover the reward function from an expert's action.
- Apprenticeship learning helps to improve the efficiency of learning.

# DeepMind in reinforcement learning: Chess, Go & RTS Games

